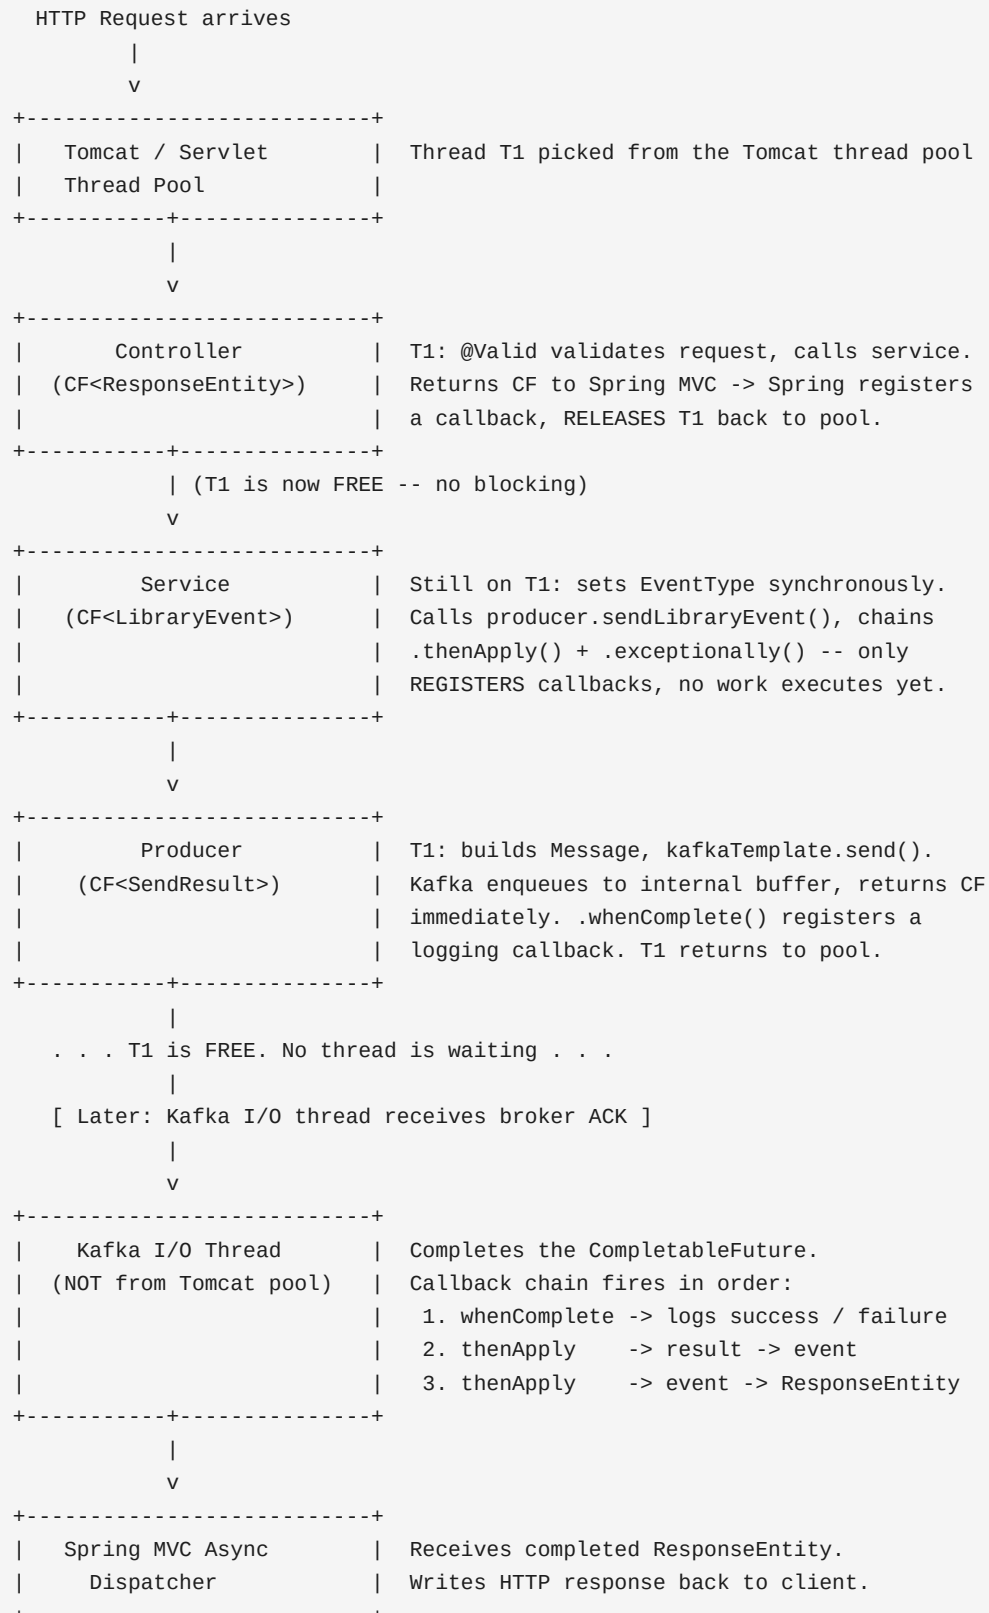


The Key Insight: Threads vs Work

In a traditional (blocking) approach, one thread does all the work AND waits for the broker ACK before returning. In this implementation, threads hand off work and return to the pool immediately -- the actual I/O wait happens on a Kafka-managed I/O thread, not on a Tomcat Servlet thread.

Request Flow -- Step by Step



Callback Chain Diagram

```
kafkaTemplate.send() -----> CF<SendResult>
                        |
                        broker ACK (Kafka I/O thread)
                        |
whenComplete -----> logs partition + offset      [Producer]
                        |
thenApply -----> SendResult -> LibraryEvent      [Service]
                        |
thenApply -----> LibraryEvent -> ResponseEntity [Controller]
                        |
Spring MVC -----> writes HTTP 201 response

(if broker rejects / timeout)
                        |
whenComplete -----> logs error
                        |
exceptionally -----> throws LibraryEventPublishException
                        |
Spring MVC -----> @ControllerAdvice -> HTTP 500
```

Code Walkthrough: thenApply / exceptionally Chain

```
return libraryEventProducer.sendLibraryEvent(event) // (1)
    .thenApply(result -> event)                      // (2)
    .exceptionally(ex -> {                          // (3)
        throw new LibraryEventPublishException(
            "Unable to publish ADD library event",
            ex.getCause()                          // (4)
        );
    });
```

(1) `sendLibraryEvent(event)` -- returns `CF<SendResult>` immediately. No waiting. The future completes later when the Kafka I/O thread gets the broker ACK.

(2) `.thenApply(result -> event)` -- registers a transformation: when the future completes successfully, ignore the `SendResult` and return the original event object. Type changes: `CF<SendResult> -> CF<LibraryEvent>`.

(3) `.exceptionally(ex -> { throw ... })` -- registers an error callback: if the future completes exceptionally, re-wrap the cause as `LibraryEventPublishException` and keep the future failed so `@ControllerAdvice` can handle it.

(4) `ex.getCause()` -- the `ex` inside `exceptionally` is always a `CompletionException` (Java's internal wrapper). The real error is one level deeper. Using `getCause()` avoids wrapping `CompletionException` inside `LibraryEventPublishException`.

Why This Matters Under Load

Scenario	Blocking (old)	Async (new)
Threads per request	1 held until broker ACK	~0 (released in microseconds)
100 req, 50ms Kafka latency	100 threads waiting	Near 0 threads waiting
Throughput ceiling	Thread pool size (~200)	Kafka I/O throughput
Tomcat thread usage	High -- proportional to Kafka latency	Minimal -- only for request parsing

A Tomcat default pool has ~200 threads. In the blocking approach each request held a thread for the full Kafka

Why it's End-to-End Asynchronous
In the 100% asynchronous approach, threads are returned to the pool in microseconds, allowing the same 200 threads to handle thousands of concurrent in-flight Kafka requests.